

Harry Nguyen

+1 470-667-9000 | harrynguyenswe@gmail.com | linkedin.com/in/harrynguyen26 | github.com/HarryNguyen2662 | timoto.ai

EDUCATION

Georgia State University

Expected May 2027

Bachelor of Science in Computer Science • GPA: 3.75/4.0

- Coursework: Distributed Systems, Operating Systems, Database Systems, Computer Networks, Data Structures
- Honors: Presidential Scholarship, Anita Andrew Fellowship, Roxie Alexander Memorial Award

WORK EXPERIENCE

Google | Software Engineer Intern — Chrome Browser

May 2025 – Aug 2025

- Architected event-driven **TypeScript/React** observer system decoupling UI state propagation across **25K+ lines** of Chromium at **95% test coverage**; reviewed and adopted into production branch by senior Chrome engineers
- Designed **C++ IPC transport layer** with Protocol Buffers serving **3B+ active users** at **sub-50ms p99, 10K+ req/sec** — selected Protobuf over custom serialization for schema evolution and zero-copy cross-process message passing
- Profiled settings navigation as p99 bottleneck at 1,200ms; designed **lock-free concurrent trie search** eliminating mutex contention — **96% latency reduction**, zero production regressions in Chrome stable channel

Develop for Good | Software Engineer Intern

May 2024 – Aug 2024

- Built **Django REST BaaS** on **AWS** with JWT-based auth and auto-scaling — **500+ concurrent users, 90% deployment time reduction** via GitHub Actions CI/CD
- Diagnosed N+1 query bottleneck degrading response to 3s+; redesigned with **PostgreSQL** indexing achieving **sub-100ms** for 10,000+ records

CoderPush | Software Engineer Intern

Sep 2023 – Dec 2023

- Diagnosed **DynamoDB hot-partition failure** in high-throughput data pipeline at **9,000+ req/sec**; redesigned partition key strategy — **30% read throughput improvement** under sustained concurrency
- Designed idempotent **REST APIs** with **Redis** distributed caching (**85% cache hit rate**) and exponential backoff — enforced exactly-once semantics preventing duplicate transactions under network partitions

PERSONAL PROJECTS

TiMoto AI – Full-Stack AI Platform | *Python, Django, gRPC, TypeScript, AWS, Terraform* timoto.ai

- **Sole engineer** — built and operate end-to-end AI platform: **Python/Django** REST backend, **gRPC** inter-service layer, **React** frontend, **PostgreSQL**, and **AWS** cloud infrastructure
- Resolved **production deadlock under concurrent gRPC calls** — traced root cause via structured logging to circular resource acquisition; redesigned call sequencing — **100% success rate** at **sub-50ms p99**
- Selected **vLLM** with PagedAttention for zero OOM failures under concurrent load — continuous batching manages KV cache without page faults at production traffic
- Architected multi-AZ **ECS Fargate** with **Terraform** IaC and circuit breaker pattern — **99.9% uptime, 44% cost reduction** to \$40–60/month; auto-rollback on health check failure

Pulumi – Open Source Contributor | *Go, TypeScript, IaC*

24.4K★ github.com/pulumi/pulumi

- Contributed Go CLI features enabling multi-cloud (AWS/Azure/GCP) provisioning — under active review by core maintainers
- Analyzed **Raft/Paxos consensus** in distributed state synchronization layer — studied linearizability guarantees and correctness under concurrent operations and partial failures

TECHNICAL SKILLS

Languages: Python, TypeScript, Go, C++, JavaScript, Java, Rust, SQL, Bash

Web & Backend: Django, FastAPI, React, REST APIs, gRPC, Protocol Buffers, Node.js, PostgreSQL, Redis

Cloud & Infrastructure: AWS (ECS Fargate, EC2, S3, RDS, DynamoDB), GCP, Kubernetes, Docker, Terraform, CI/CD, GitHub Actions

Distributed Systems: Exactly-once semantics, Circuit breakers, Lock-free data structures, Raft/Paxos, Linearizability, Fault tolerance

ML & AI: vLLM, PagedAttention, Continuous Batching, PyTorch, LLM-as-a-judge evaluation