

Harry Nguyen

+1 470-667-9000 | harry.nguyen@timoto.ai | linkedin.com/in/harrynguyen26 | github.com/HarryNguyen2662 | timoto.ai
| Atlanta, GA

PROFESSIONAL SUMMARY

CS @ Georgia State (GPA 3.75, May 2027). Primary engineer for a production GenAI serving platform at TiMoto AI -- designed and operated vLLM/PagedAttention inference (sub-50ms p99, 99.9% uptime, zero OOM failures), built LLM evaluation pipelines, and owned multi-AZ AWS infrastructure with Terraform. At Google I shipped C++ IPC to Chrome stable for 3B+ users. I build production LLM systems that hit SLOs and hold up under concurrency -- ready to own customer-facing AI deployments at Databricks scale. Case study: [timoto.ai](#)

CORE COMPETENCIES

Production LLM Serving (vLLM)

RAG & GenAI Engineering

LLM Evaluation Pipelines

MLOps / AWS ML Deployment

gRPC & Distributed Systems

LangChain / PyTorch / LoRA

Terraform IaC / ECS Fargate

Observability & Incident Response

C++ / Python / Go / TypeScript

Technical Communication

EXPERIENCE

TiMoto AI Sep 2025 -- Present

Software Engineer -- Backend, Infrastructure & ML Serving

- Primary engineer for an end-to-end **production GenAI serving platform** -- designed, deployed, and operated LLM inference (vLLM/PagedAttention) for real-time user requests on a 3-person team.
- Built **LLM evaluation pipeline using LLM-as-a-judge** -- achieved **100% evaluation success rate at sub-50ms p99**; caught 3 regressions before production promotion.
- Architected **vLLM inference engine with PagedAttention** -- eliminated KV cache memory fragmentation under concurrent load; deployed with continuous batching achieving **zero OOM failures** at production traffic.
- Debugged and fixed a production **gRPC distributed deadlock** under concurrent evaluation calls -- traced cross-service lock acquisition conflicts and redesigned call sequencing; evaluation success rate restored to 100%.
- Architected **multi-AZ ECS Fargate with Terraform IaC**, CloudWatch observability, and circuit breaker -- **99.9% uptime, 44% cost reduction** (\$40-60/month); auto-rollback on health check failure.

Google May 2025 -- Aug 2025

Software Engineer Intern -- Chrome Browser

- Delivered design documents and code reviews adopted by senior Chrome engineers; shipped **C++ IPC transport layer with Protocol Buffers** to Chrome stable serving **3B+ active users** at sub-50ms p99, 10K+ req/sec.
- Identified settings navigation as p99 bottleneck at 1,200ms; designed and implemented **lock-free concurrent trie search** eliminating mutex contention -- **96% latency reduction**, zero production regressions.
- Architected event-driven TypeScript/React system across **25K+ lines of Chromium** at **95% test coverage**; accelerated feature delivery 68%.

Software Engineer Intern

- Designed stateless BaaS on AWS -- auto-scaling policy for **500+ concurrent users**, **90% deployment time reduction** via CI/CD; explained architecture decisions (JWT vs session auth) to non-technical program stakeholders.
- Diagnosed N+1 query bottleneck (3s+ on large datasets); redesigned with PostgreSQL indexing achieving **sub-100ms** for 10,000+ records.

PROJECTS

Pulumi -- Open Source Contributor [Go / TypeScript / IaC](#)

Submitted Go CLI features and bug fixes for multi-cloud (AWS/Azure/GCP) provisioning -- under active review by core maintainers (24K+ star project). Studied Raft/Paxos consensus in distributed state synchronization layer; analyzed linearizability guarantees under partial failures.

EDUCATION

B.S. Computer Science -- Georgia State University

Expected May 2027

GPA: 3.75/4.0 | Coursework: Distributed Systems, Operating Systems, Database Systems, Computer Networks, Data Structures
Honors: Presidential Scholarship, Anita Andrew Fellowship, Roxie Alexander Memorial Award

TECHNICAL SKILLS

GenAI & ML Infra: vLLM, PagedAttention, Continuous Batching, RAG, LangChain, HuggingFace (familiar), DSPy (familiar), LoRA/QLoRA, PyTorch, LLM-as-a-judge evaluation, fine-tuning

Distributed Systems: gRPC, Protocol Buffers, Circuit breakers, Exactly-once semantics, Fault tolerance, Kafka, Multi-AZ, Event-driven architecture

Cloud & MLOps: AWS (ECS Fargate, EKS, EC2, S3, RDS), Terraform, Kubernetes, Docker, CI/CD, CloudWatch, Prometheus, Grafana, GCP

Languages: Python, C++, Go, TypeScript, Java, JavaScript, SQL, Bash, Rust

Frameworks & Databases: Django, FastAPI, Node.js, React, REST APIs, PostgreSQL, Redis, MongoDB

Learning: Apache Spark, Databricks Lakehouse Platform, Delta Lake