

Harry Nguyen

+1 470-667-9000 | nguyenharry2662@gmail.com | linkedin.com/in/harrynguyen26 | github.com/HarryNguyen2662

EDUCATION

Georgia State University

Expected May 2027

Bachelor of Science in Computer Science • **GPA: 3.75/4.0**

- Coursework: Distributed Systems, Operating Systems, Database Systems, Computer Networks, Data Structures
- Honors: Presidential Scholarship, Anita Andrew Fellowship, Roxie Alexander Memorial Award

EXPERIENCE

TiMoto AI | Software Engineer

Sep 2025 – Present

- Built full-stack AI evaluation platform as primary engineer on 3-person team – owned REST API layer (Django), cloud infra (AWS/Terraform), and ML serving end-to-end
- Architected **vLLM inference engine with PagedAttention** for LLM serving – eliminated KV cache memory fragmentation under concurrent load; deployed with continuous batching achieving **zero OOM failures** at production traffic, **sub-50ms p99**
- Designed and maintained **gRPC inter-service layer**; resolved production deadlock under concurrent calls by tracing shared resource acquisition conflicts and redesigning call sequencing – **100% evaluation success rate**
- Architected multi-AZ ECS Fargate with **Terraform IaC**, CloudWatch observability, and circuit breaker pattern – **99.9% uptime, 44% cost reduction** (\$40–60/month); auto-rollback on health check failure
- Maintained distributed production systems across REST API, gRPC, PostgreSQL, and ML serving layers – on-call rotation, runbook authoring, post-mortem driven improvements

Google | Software Engineer Intern – Chrome Browser

May 2025 – Aug 2025

- Architected event-driven **TypeScript/React system** with observer pattern decoupling UI state propagation – accelerated feature delivery **68%** across **25K+ lines** of Chromium at **95% test coverage**
- Designed **C++ IPC transport layer** with Protocol Buffers – shipped to Chrome stable channel serving **3B+ active users** at **sub-50ms p99, 10K+ req/sec**
- Identified settings navigation as p99 bottleneck at 1,200ms; designed and implemented **lock-free concurrent trie search** eliminating mutex contention – **96% latency reduction**, zero regressions
- Delivered design documents and code reviews at **95% test coverage**; changes adopted into production branch reviewed by senior Chrome engineers

Develop for Good | Software Engineer Intern

May 2024 – Aug 2024

- Built stateless BaaS on AWS – JWT auth for horizontal scalability; auto-scaling policy supporting **500+ concurrent users, 90% deployment time reduction** via CI/CD (GitHub Actions)
- Diagnosed N+1 query bottleneck degrading response to 3s+ on large datasets; redesigned with PostgreSQL indexing achieving **sub-100ms** for 10,000+ records

PROJECTS

Pulumi – Contributor | *Go, TypeScript, IaC*

github.com/pulumi

- Submitted Go CLI features and bug fixes enabling multi-cloud (AWS/Azure/GCP) provisioning – under active review by core maintainers
- Analyzed **Raft/Paxos consensus** in distributed state synchronization layer – studied linearizability guarantees under concurrent operations and partial failures

SKILLS

Languages: Java, Python, TypeScript, Go, C++, JavaScript, SQL, Bash, Rust

Full-Stack & APIs: React, Django, FastAPI, Node.js, Spring Boot, REST APIs, gRPC, Protocol Buffers

Cloud & Infrastructure: AWS (ECS Fargate, Lambda, EC2, S3, RDS), Terraform, K8s, Docker, CI/CD, CloudWatch, Prometheus, Grafana

ML & AI: vLLM, PagedAttention, Continuous Batching, LoRA/QLoRA, PyTorch, LangChain, LLM-as-a-judge evaluation

Databases: PostgreSQL, Redis, MongoDB, DynamoDB

AI Dev Tools: Claude Code, GitHub Copilot, Codex, Cursor