

Harry Nguyen

+1 470-667-9000 | nguyenharry2662@gmail.com | linkedin.com/in/harrynguyen26 | github.com/HarryNguyen2662

EDUCATION

Georgia State University

Expected May 2027

Bachelor of Science in Computer Science • **GPA: 3.75/4.0**

- Coursework: Distributed Systems, Operating Systems, Database Systems, Computer Networks, Data Structures
- Honors: Presidential Scholarship, Anita Andrew Fellowship, Roxie Alexander Memorial Award

EXPERIENCE

TiMoto AI | Software Engineer – Backend, Infrastructure & ML Serving

Sep 2025 – Present

- Built and operated **distributed AI agent systems end-to-end** as primary engineer on a 3-person team – owned backend, cloud infra, and ML serving from design to production in a **continuous-deploy environment**
- Designed and maintained **gRPC inter-service layer**; resolved production deadlock under concurrent calls by tracing shared resource acquisition conflicts and redesigning call sequencing – **100% evaluation success rate** at **sub-50ms p99**
- Architected **vLLM inference engine with PagedAttention** for LLM serving – eliminated KV cache memory fragmentation under concurrent load; deployed with continuous batching achieving **zero OOM failures** at production traffic
- Architected multi-AZ ECS Fargate with **Terraform IaC**, CloudWatch + Prometheus observability, and circuit breaker – **99.9% uptime**, **44% cost reduction** (\$40–60/month); auto-rollback on health check failure
- Triaged and resolved production incidents across gRPC, PostgreSQL, and ML serving layers – documented runbooks, applied post-mortem learnings, and participated in on-call rotation

Google | Software Engineer Intern – Chrome Browser

May 2025 – Aug 2025

- Designed **C++ IPC transport layer** with Protocol Buffers – selected over custom serialization for schema evolution; shipped to Chrome stable serving **3B+ active users** at **sub-50ms p99**, **10K+ req/sec**
- Identified settings navigation as p99 bottleneck at 1,200ms; implemented **lock-free concurrent trie** eliminating mutex contention – **96% latency reduction**, zero production regressions
- Architected event-driven TypeScript/React system with observer pattern; collaborated with product, design, and infra teams from prototype to production – **68% faster feature delivery** across **25K+ lines** of Chromium at **95% test coverage**

Develop for Good | Software Engineer Intern

May 2024 – Aug 2024

- Designed stateless BaaS on AWS – JWT over session auth for horizontal scalability; **90% deploy time reduction** via CI/CD (GitHub Actions); auto-scaling supporting **500+ concurrent users**
- Diagnosed N+1 bottleneck degrading response to 3s+ on large datasets; redesigned with PostgreSQL indexing achieving **sub-100ms** for 10,000+ records

PROJECTS

Pulumi – Contributor | *Go, TypeScript, IaC*

github.com/pulumi

- Submitted Go CLI features and bug fixes enabling multi-cloud (AWS/Azure/GCP) provisioning – under active review by core maintainers
- Analyzed **Raft/Paxos consensus** in distributed state synchronization layer – studied linearizability guarantees and correctness under concurrent operations and partial failures

SKILLS

Distributed Systems: gRPC, Protocol Buffers, Circuit breakers, Exactly-once semantics, Fault tolerance, Kafka, Multi-AZ, Event-driven architecture

ML & AI Infrastructure: vLLM, PagedAttention, Continuous Batching, LoRA/QLoRA, PyTorch, LangChain, LLM-as-a-judge evaluation, AI agent orchestration

Languages: C++, Python, Go, TypeScript, Java, JavaScript, SQL, Bash, Rust

Cloud & Infrastructure: AWS (ECS Fargate, EKS, EC2, S3, RDS), Terraform, K8s, Docker, CI/CD, CloudWatch, Prometheus, Grafana, GCP

Frameworks & Databases: Django, FastAPI, Node.js, React, REST APIs, PostgreSQL, Redis, MongoDB

AI Dev Tools: Claude Code, GitHub Copilot, Codex, Cursor