

Harry Nguyen

+1 470-667-9000 | nguyenharry2662@gmail.com | linkedin.com/in/harrynguyen26 | github.com/HarryNguyen2662

EDUCATION

Georgia State University

Expected May 2027

Bachelor of Science in Computer Science • **GPA: 3.75/4.0**

- Coursework: Distributed Systems, Operating Systems, Database Systems, Computer Networks, Data Structures
- Honors: Presidential Scholarship, Anita Andrew Fellowship, Roxie Alexander Memorial Award

EXPERIENCE

TiMoto AI | Software Engineer

Sep 2025 – Present

- Designed and deployed **LLM-as-a-judge evaluation pipeline** and **vLLM inference engine with PagedAttention** – achieved **100% evaluation success rate** at **sub-50ms p99**; integrated **LangChain** for LLM orchestration in production
- Built and owned full-stack backend: **Python (Django/FastAPI)** REST APIs, **PostgreSQL** schema design, **Redis** caching layer, and **gRPC** inter-service communication – resolved production deadlock under concurrent calls
- Architected multi-AZ ECS Fargate with **Terraform IaC** and **CloudWatch observability** – **99.9% uptime**, **44% cost reduction** (\$40–60/month); circuit breaker pattern with auto-rollback on health check failure
- Led backend, cloud infrastructure, and ML serving for a 3-person engineering team – designed, deployed, and operated **distributed production systems** end-to-end

Google | Software Engineer Intern – Chrome Browser

May 2025 – Aug 2025

- Architected event-driven **TypeScript/React** system with observer pattern decoupling UI state propagation – **68% feature delivery acceleration** across **25K+ lines** of Chromium at **95% test coverage**
- Designed **C++ IPC transport layer** with Protocol Buffers – shipped to Chrome stable channel serving **3B+ active users** at **sub-50ms p99**, **10K+ req/sec**
- Implemented **lock-free concurrent trie search** eliminating mutex contention – **96% latency reduction**, zero production regressions
- Delivered design documents and code reviews with Chrome infrastructure team; changes adopted into production branch by senior Chrome engineers

Develop for Good | Software Engineer Intern

May 2024 – Aug 2024

- Designed stateless BaaS on AWS – auto-scaling supporting **500+ concurrent users**, **90% deployment time reduction** via CI/CD (GitHub Actions)
- Diagnosed N+1 query bottleneck; redesigned with **PostgreSQL indexing** achieving **sub-100ms** response for 10,000+ records

PROJECTS

Pulumi – Contributor | Go, TypeScript, IaC

github.com/pulumi

- Submitted Go CLI features and bug fixes enabling multi-cloud (AWS/Azure/GCP) provisioning – under active review by core maintainers
- Analyzed **Raft/Paxos consensus** in distributed state synchronization layer – studied linearizability guarantees and correctness under concurrent operations and partial failures

SKILLS

ML & AI Infrastructure: vLLM, PagedAttention, Continuous Batching, LLM-as-a-judge evaluation, LangChain, LoRA/QLoRA, PyTorch

Languages: Python, TypeScript, Go, C++, Java, JavaScript, SQL, Bash, Rust

Frameworks & Databases: FastAPI, Django, Node.js, React, REST APIs, PostgreSQL, Redis, MongoDB

Distributed Systems: gRPC, Protocol Buffers, Circuit breakers, Exactly-once semantics, Fault tolerance, Kafka, Multi-AZ, Event-driven architecture

Cloud & Infrastructure: AWS (ECS Fargate, EKS, EC2, S3, RDS), Terraform, K8s, Docker, CI/CD, CloudWatch, Prometheus, Grafana

AI Dev Tools: Claude Code, GitHub Copilot, Codex, Cursor