

Harry Nguyen

+1 470-667-9000 | nguyenharry2662@gmail.com | linkedin.com/in/harrynguyen26 | github.com/HarryNguyen2662

EDUCATION

Georgia State University

Expected May 2027

Bachelor of Science in Computer Science • GPA: 3.75/4.0

- Coursework: Distributed Systems, Operating Systems, Database Systems, Computer Networks, Data Structures
- Honors: Presidential Scholarship, Anita Andrew Fellowship, Roxie Alexander Memorial Award

EXPERIENCE

TiMoto AI | Software Engineer

Sep 2025 – Present

- Architected **LangChain-powered AI agent serving** and **LLM-as-a-judge evaluation pipeline** – deployed in production with **100% evaluation success rate** at **sub-50ms p99**; integrated **LLM APIs (Claude, GPT)** for multi-step agent orchestration
- Led backend, cloud infrastructure, and ML serving for a 3-person engineering team – designed and deployed **distributed production systems** end-to-end using **vLLM with PagedAttention**; zero OOM failures at production traffic
- Designed and maintained **gRPC inter-service layer**; resolved production deadlock under concurrent calls by tracing shared resource acquisition conflicts and redesigning call sequencing
- Architected multi-AZ ECS Fargate with **Terraform IaC**, CloudWatch observability, and circuit breaker pattern – **99.9% uptime**, **44% cost reduction** (\$40–60/month); on-call rotation with runbooks and post-mortem process

Google | Software Engineer Intern – Chrome Browser

May 2025 – Aug 2025

- Collaborated with Chrome infrastructure team to deliver **design documents** and code reviews – explained Protocol Buffers vs custom serialization tradeoffs to senior engineers; changes adopted into production branch serving **3B+ active users**
- Designed **C++ IPC transport layer** with Protocol Buffers – shipped to Chrome stable channel at **sub-50ms p99**, **10K+ req/sec**
- Identified settings navigation as p99 bottleneck at 1,200ms; designed and implemented **lock-free concurrent trie search** eliminating mutex contention – **96% latency reduction**, zero production regressions
- Architected event-driven TypeScript/React system with observer pattern decoupling UI state propagation – **68% feature delivery acceleration** across **25K+ lines** of Chromium at **95% test coverage**

Develop for Good | Software Engineer Intern

May 2024 – Aug 2024

- Designed stateless BaaS on AWS – selected JWT over session auth for horizontal scalability; auto-scaling policy supporting **500+ concurrent users**, **90% deployment time reduction** via CI/CD (GitHub Actions)
- Diagnosed N+1 query bottleneck degrading response to 3s+ on large datasets; redesigned with **PostgreSQL indexing** achieving **sub-100ms** for 10,000+ records

PROJECTS

Pulumi – Contributor | Go, TypeScript, IaC

github.com/pulumi

- Submitted Go CLI features and bug fixes enabling multi-cloud (AWS/Azure/GCP) provisioning – under active review by core maintainers
- Analyzed **Raft/Paxos consensus** in distributed state synchronization layer – studied linearizability guarantees and correctness under concurrent operations and partial failures

SKILLS

ML & AI Infrastructure: LangChain, LangGraph, vLLM, PagedAttention, Continuous Batching, LLM-as-a-judge evaluation, LLM APIs (Claude, GPT), LoRA/QLoRA, PyTorch

Languages: Python, TypeScript, Go, JavaScript, C++, Java, SQL, Bash, Rust

Frameworks & Databases: FastAPI, Django, Node.js, React, REST APIs, PostgreSQL, Redis, MongoDB

Cloud & Infrastructure: AWS (ECS Fargate, EKS, EC2, S3, RDS), Terraform, Kubernetes, Docker, CI/CD, CloudWatch, Prometheus, Grafana, GCP

Distributed Systems: gRPC, Protocol Buffers, Circuit breakers, Exactly-once semantics, Fault tolerance, Multi-AZ, Event-driven architecture

AI Dev Tools: Claude Code, GitHub Copilot, Codex, Cursor