

Harry Nguyen

+1 470-667-9000 | nguyenharry2662@gmail.com | linkedin.com/in/harrynguyen26 | github.com/HarryNguyen2662

EDUCATION

Georgia State University

Expected May 2027

Bachelor of Science in Computer Science • GPA: 3.75/4.0

- Coursework: Distributed Systems, Operating Systems, Database Systems, Computer Networks, Data Structures
- Honors: Presidential Scholarship, Anita Andrew Fellowship, Roxie Alexander Memorial Award

EXPERIENCE

TiMoto AI | Software Engineer

Sep 2025 – Present

- Led backend, ML serving, and product API development for a 3-person engineering team – built **Python (FastAPI/Django) REST APIs** and **LLM integrations (Claude, GPT)** with **LangChain** orchestration deployed to production
- Architected **vLLM inference engine with PagedAttention** for AI model serving – eliminated KV cache memory fragmentation under concurrent load; continuous batching achieving **zero OOM failures** at production traffic with **sub-50ms p99**
- Designed and maintained **gRPC inter-service layer**; resolved production deadlock by tracing shared resource acquisition conflicts and redesigning call sequencing – **100% evaluation success rate**
- Architected multi-AZ ECS Fargate with **Terraform IaC**, CloudWatch observability, and circuit breaker pattern – **99.9% uptime, 44% cost reduction** (\$40–60/month); auto-rollback on health check failure

Google | Software Engineer Intern – Chrome Browser

May 2025 – Aug 2025

- Architected **event-driven TypeScript/React system** with observer pattern decoupling UI state propagation – **68% feature delivery acceleration** across **25K+ lines** of Chromium at **95% test coverage**; shipped to **3B+ active users**
- Designed **C++ IPC transport layer** with Protocol Buffers – selected Protocol Buffers over custom serialization for schema evolution and cross-language compatibility; shipped to Chrome stable channel at **sub-50ms p99, 10K+ req/sec**
- Identified settings navigation as p99 bottleneck at 1,200ms; designed and implemented **lock-free concurrent trie search** eliminating mutex contention – **96% latency reduction**, zero production regressions
- Collaborated with Chrome infrastructure team to deliver design documents and code reviews; changes adopted into production branch reviewed by senior Chrome engineers

Develop for Good | Software Engineer Intern

May 2024 – Aug 2024

- Designed stateless BaaS on AWS – selected JWT over session auth for horizontal scalability; auto-scaling policy supporting **500+ concurrent users, 90% deployment time reduction** via CI/CD (GitHub Actions)
- Diagnosed N+1 query bottleneck degrading response to 3s+ on large datasets; redesigned with **PostgreSQL indexing** achieving **sub-100ms** for 10,000+ records

PROJECTS

Pulumi – Contributor | Go, TypeScript, IaC

github.com/pulumi

- Submitted Go CLI features and bug fixes enabling multi-cloud (AWS/Azure/GCP) provisioning – under active review by core maintainers
- Analyzed **Raft/Paxos consensus** in distributed state synchronization layer – studied linearizability guarantees and correctness under concurrent operations and partial failures

SKILLS

ML & AI Infrastructure: vLLM, PagedAttention, Continuous Batching, LLM APIs (Claude, GPT), LangChain, LLM-as-a-judge evaluation, LoRA/QLoRA, PyTorch

Languages: Python, TypeScript, Go, JavaScript, C++, Java, SQL, Bash, Rust

Frameworks & Databases: FastAPI, Django, Node.js, React, REST APIs, PostgreSQL, Redis, MongoDB

Cloud & Infrastructure: AWS (ECS Fargate, EKS, EC2, S3, RDS), Terraform, Kubernetes, Docker, CI/CD, CloudWatch, Prometheus, Grafana

Distributed Systems: gRPC, Protocol Buffers, Circuit breakers, Exactly-once semantics, Fault tolerance, Multi-AZ, Event-driven architecture

AI Dev Tools: Claude Code, GitHub Copilot, Codex, Cursor