

Harry Nguyen

+1 470-667-9000 | nguyenharry2662@gmail.com | linkedin.com/in/harrynguyen26 | github.com/HarryNguyen2662

EDUCATION

Georgia State University

Expected May 2027

Bachelor of Science in Computer Science • GPA: 3.75/4.0

- Coursework: Distributed Systems, Operating Systems, Database Systems, Computer Networks, Data Structures
- Honors: Presidential Scholarship, Anita Andrew Fellowship, Roxie Alexander Memorial Award

EXPERIENCE

TiMoto AI | Software Engineer

Sep 2025 – Present

- Deployed **vLLM with PagedAttention** for production LLM inference – solved KV cache memory fragmentation under concurrent load; **zero OOM failures** at production traffic; **continuous batching** throughput optimization; benchmarked vLLM vs. naive HF inference on throughput, p99 latency, and memory utilization
- Built **LLM-as-a-judge evaluation pipeline** – AI infrastructure for automated regression detection before deployment; **LangChain** orchestration + LLM API integration (Claude, GPT); **PyTorch** ecosystem ML serving; **sub-50ms p99** inference SLO
- Primary engineer on 3-person team – designed **gRPC inter-service layer** (exactly-once delivery, circuit breaker auto-rollback); multi-AZ **AWS ECS Fargate** with **Terraform IaC**; **99.9% uptime SLO**; **44% infrastructure cost reduction** (\$40–60/mo)

Google | Software Engineer Intern – Chrome Browser

May 2025 – Aug 2025

- Designed **C++ IPC transport layer** with **Protocol Buffers** – evaluated Protobuf vs. custom serialization tradeoffs (schema evolution, memory overhead, cross-language compatibility); shipped to Chrome stable serving **3B+ active users** at **sub-50ms p99**
- Diagnosed p99 bottleneck at 1,200ms; profiled critical path; implemented **lock-free concurrent trie search** with **CAS operations** eliminating mutex contention – **96% latency reduction**, zero production regressions; production C++ **performance optimization** at scale
- Architected event-driven **TypeScript/React** system with observer pattern – **68% feature delivery acceleration** across **25K+ lines** of Chromium at **95% test coverage**; technical documentation + code reviews adopted by senior Chrome engineers

Develop for Good | Software Engineer Intern

May 2024 – Aug 2024

- Designed stateless BaaS on **AWS** – JWT over session auth for horizontal scalability; auto-scaling supporting **500+ concurrent users**; **90% deployment time reduction** via CI/CD (GitHub Actions)
- Diagnosed N+1 query bottleneck degrading response to 3s+; redesigned with **PostgreSQL composite indexing** achieving **sub-100ms** for 10,000+ records

PROJECTS

Pulumi – Contributor | Go, TypeScript, Infrastructure-as-Code

github.com/pulumi

- Submitted **Go CLI features** for multi-cloud (AWS/Azure/GCP) provisioning – open-source developer tooling under active review by core maintainers
- Analyzed **Raft/Paxos consensus** in distributed state synchronization – studied linearizability guarantees and correctness under concurrent operations and partial failures

SKILLS

ML & AI Infrastructure: vLLM, PagedAttention, Continuous Batching, PyTorch, LangChain, LLM-as-a-judge evaluation, LLM APIs (Claude, GPT), LoRA/QLoRA

Distributed Systems: gRPC, Protocol Buffers, Exactly-once semantics, Circuit breakers, Multi-AZ, Fault tolerance, Event-driven architecture

Languages: C++, Python, Go, TypeScript, Java, JavaScript, SQL, Bash, Rust

Cloud & Infrastructure: AWS (ECS Fargate, EKS, EC2, S3, RDS), Terraform, Kubernetes, Docker, CI/CD (GitHub Actions), CloudWatch, Prometheus, Grafana

Frameworks & Databases: FastAPI, Django, Node.js, React, REST APIs, PostgreSQL, Redis, MongoDB

AI Dev Tools: Claude Code, GitHub Copilot, Codex, Cursor