

Harry Nguyen

+1 470-667-9000 | nguyenharry2662@gmail.com | linkedin.com/in/harrynguyen26 | github.com/HarryNguyen2662

EDUCATION

Georgia State University

Expected May 2027

Bachelor of Science in Computer Science • **GPA: 3.75/4.0**

- Coursework: Distributed Systems, Operating Systems, Database Systems, Computer Networks, Data Structures
- Honors: Presidential Scholarship, Anita Andrew Fellowship, Roxie Alexander Memorial Award

EXPERIENCE

TiMoto AI | Software Engineer

Sep 2025 – Present

- Built **LLM-as-a-judge AI evaluation pipeline** – automated deployment decisions using **LLM APIs**; **LangChain** orchestration + Claude/GPT integration; quality gates before deployment; replaced manual review with continuous automated enforcement; **zero production regressions**; primary engineer on 3-person team
- Deployed **vLLM with PagedAttention** for production **ML model serving** – solved KV cache memory fragmentation under concurrent inference load; continuous batching throughput optimization; **zero OOM failures**; monitored with CloudWatch + Prometheus/Grafana
- Designed **gRPC inter-service data pipeline** with **exactly-once semantics** – traced production deadlock under concurrent load; redesigned call sequencing; **100% evaluation success rate** at **sub-50ms p99**; **FastAPI** + **Django** REST layer serving external API clients
- Architected **multi-AZ AWS ECS Fargate** with **Terraform IaC**, circuit breaker pattern, and auto-rollback – **99.9% uptime SLO**; on-call rotation, root-cause analysis, runbooks, post-mortem; **44% infrastructure cost reduction** (\$40–60/mo)

Google | Software Engineer Intern – Chrome Browser

May 2025 – Aug 2025

- Designed **C++ IPC transport layer** with **Protocol Buffers** at Google's security review standard – evaluated Protobuf vs. custom serialization (schema evolution, backwards-compatibility, cross-language support); shipped to Chrome stable serving **3B+ active users** at **sub-50ms p99**; **95% test coverage**; zero production regressions
- Profiled settings navigation at 1,200ms p99; implemented **lock-free concurrent trie search** with **CAS operations** eliminating mutex contention – **96% latency reduction**, zero regressions; design adopted by senior Chrome engineers
- Architected event-driven **TypeScript/React** system with observer pattern – **68% feature delivery acceleration** across **25K+ lines** of Chromium at **95% test coverage**

Develop for Good | Software Engineer Intern

May 2024 – Aug 2024

- Designed stateless **BaaS on AWS** – selected **JWT** over session auth for stateless horizontal scalability; auto-scaling supporting **500+ concurrent users**; **90% deployment time reduction** via CI/CD (**GitHub Actions**)
- Diagnosed N+1 query bottleneck degrading API response to 3s+ on large datasets; redesigned with **PostgreSQL composite indexing** achieving **sub-100ms** for 10,000+ records

PROJECTS

Pulumi – Contributor | *Go, TypeScript, Multi-Cloud IaC*

github.com/pulumi

- Submitted **Go CLI features** for multi-cloud (AWS/Azure/GCP) provisioning – open-source developer tooling under active review by core maintainers; 24.4K+ star project
- Analyzed **Raft/Paxos consensus** in distributed state synchronization – studied linearizability guarantees and correctness under concurrent operations and partial failures

SKILLS

ML & AI Infrastructure: vLLM, PagedAttention, LLM-as-a-judge evaluation, Continuous Batching, LangChain, LLM APIs (Claude, GPT), PyTorch

Languages: Python, Go, C++, TypeScript, JavaScript, Java, Rust, SQL, Bash

Distributed Systems: gRPC, Protocol Buffers, Exactly-once semantics, Circuit breakers, Multi-AZ, Fault tolerance, Event-driven architecture

Cloud & Infrastructure: AWS (ECS Fargate, EKS, EC2, S3, RDS), Terraform, Kubernetes, Docker, CI/CD (GitHub Actions), CloudWatch, Prometheus, Grafana, GCP

Frameworks & APIs: FastAPI, Django, React, Node.js, REST APIs, PostgreSQL, Redis, MongoDB, HTML/CSS

AI Dev Tools: Claude Code, GitHub Copilot, Codex, Cursor