

Harry Nguyen

+1 470-667-9000 | nguyenharry2662@gmail.com | linkedin.com/in/harrynguyen26 | github.com/HarryNguyen2662

EDUCATION

Georgia State University

Expected May 2027

Bachelor of Science in Computer Science • **GPA: 3.75/4.0**

- Coursework: Distributed Systems, Operating Systems, Database Systems, Computer Networks, Data Structures
- Honors: Presidential Scholarship, Anita Andrew Fellowship, Roxie Alexander Memorial Award
- Activities: Amazon University Event – BEN Tech Cloud Connect ATL

EXPERIENCE

TiMoto AI | Software Engineer

Sep 2025 – Present

- Architected **vLLM inference engine with PagedAttention** for LLM serving — eliminated KV cache memory fragmentation under concurrent load; deployed with continuous batching achieving **zero OOM failures** at production traffic; primary ML infra engineer on 3-person team
- Built **LLM-as-a-judge evaluation pipeline** — automated model quality scoring with calibrated criteria, surfaced regressions in CI loop, and drove **model optimization** and **data processing** workflows end-to-end
- Designed and maintained **gRPC inter-service layer**; resolved production deadlock under concurrent calls by tracing shared resource conflicts — **100% evaluation success rate** at **sub-50ms p99**
- Architected multi-AZ ECS Fargate with **Terraform IaC**, CloudWatch observability, and circuit breaker pattern — **99.9% uptime**, **44% cost reduction** (\$40–60/month); auto-rollback on health check failure
- Triaged and resolved production incidents across gRPC, PostgreSQL, and ML serving layers via root-cause analysis; documented runbooks and applied post-mortem learnings; participated in on-call rotation

Google | Software Engineer Intern — Chrome Browser

May 2025 – Aug 2025

- Designed **C++ IPC transport layer** with Protocol Buffers — shipped to Chrome stable channel serving **3B+ active users** at **sub-50ms p99**, **10K+ req/sec**; code reviewed by senior Chrome engineers at **95% test coverage**
- Identified settings navigation as p99 bottleneck at 1,200ms; implemented **lock-free concurrent trie search** eliminating mutex contention — **96% latency reduction**, zero production regressions
- Architected event-driven TypeScript/React system with observer pattern — accelerated feature delivery **68%** across **25K+ lines** of Chromium at **95% test coverage**
- Collaborated with Chrome infrastructure team on design documents and code reviews; changes adopted into production branch

Develop for Good | Software Engineer Intern

May 2024 – Aug 2024

- Designed stateless BaaS on AWS — JWT auth for horizontal scalability; auto-scaling supporting **500+ concurrent users**, **90% deployment time reduction** via CI/CD (GitHub Actions)
- Diagnosed N+1 query bottleneck at 3s+; redesigned with PostgreSQL indexing and **data processing** optimizations achieving **sub-100ms** for 10,000+ records

PROJECTS

Pulumi – Contributor | Go, TypeScript, Infrastructure-as-Code (IaC)

github.com/pulumi

- Submitted Go CLI features and bug fixes enabling multi-cloud (AWS/Azure/GCP) provisioning — under active review by core maintainers
- Analyzed **Raft/Paxos consensus** in distributed state synchronization layer — studied linearizability guarantees and correctness under concurrent operations and partial failures

SKILLS

ML & AI Infrastructure: vLLM, PagedAttention, Continuous Batching, LoRA/QLoRA, PyTorch, LangChain, LLM-as-a-judge evaluation, model optimization

Languages: Python, C++, Go, TypeScript, Java, JavaScript, SQL, Bash, Rust

Distributed Systems: gRPC, Protocol Buffers, Circuit breakers, Exactly-once semantics, Fault tolerance, Kafka, Multi-AZ, Event-driven architecture

Cloud & Infrastructure: AWS (ECS Fargate, EKS, EC2, S3, RDS), Terraform, Kubernetes, Docker, CI/CD, CloudWatch, Prometheus, Grafana, GCP

Frameworks & Databases: FastAPI, Django, Node.js, React, REST APIs, PostgreSQL, Redis, MongoDB

AI Dev Tools: Claude Code, GitHub Copilot, Codex, Cursor