

Harry Nguyen

+1 470-667-9000 | nguyenharry2662@gmail.com | linkedin.com/in/harrynguyen26 | github.com/HarryNguyen2662

EDUCATION

Georgia State University

Expected May 2027

Bachelor of Science in Computer Science • GPA: 3.75/4.0

- Coursework: Distributed Systems, Operating Systems, Database Systems, Computer Networks, Data Structures
- Honors: Presidential Scholarship, Anita Andrew Fellowship, Roxie Alexander Memorial Award

EXPERIENCE

TiMoto AI | Software Engineer

Sep 2025 – Present

- Built **LLM-as-a-judge evaluation pipeline** with LangChain – defined scoring rubrics for factual and temporal accuracy, integrated with inference layer achieving **100% evaluation success rate at sub-50ms p99**; directly analogous to freshness quality evaluation for LLM outputs
- Architected **vLLM inference engine with PagedAttention** for real-time LLM serving – eliminated KV cache memory fragmentation; deployed with continuous batching achieving **zero OOM failures** at production traffic
- Led backend, cloud infrastructure, and ML serving as primary engineer on 3-person team – designed, deployed, and operated **distributed production systems** end-to-end
- Designed and maintained **gRPC inter-service layer**; resolved production deadlock by tracing shared resource acquisition conflicts – **100% uptime** maintained at **sub-50ms p99**
- Architected multi-AZ ECS Fargate with **Terraform IaC** and CloudWatch observability – **99.9% uptime, 44% cost reduction** (\$40–60/month)

Google | Software Engineer Intern – Chrome Browser

May 2025 – Aug 2025

- Designed **C++ IPC transport layer** with Protocol Buffers – selected Protocol Buffers for schema evolution and cross-language compatibility; shipped to Chrome stable channel serving **3B+ active users** at **sub-50ms p99, 10K+ req/sec**
- Identified settings navigation as p99 bottleneck at 1,200ms; implemented **lock-free concurrent trie search** in C++ eliminating mutex contention – **96% latency reduction**, zero production regressions
- Architected event-driven TypeScript/React system – accelerated feature delivery **68%** across **25K+ lines** of Chromium at **95% test coverage**
- Collaborated with Chrome infrastructure team; design documents and code reviews adopted into production branch by senior Chrome engineers

Develop for Good | Software Engineer Intern

May 2024 – Aug 2024

- Designed stateless BaaS on AWS supporting **500+ concurrent users**; **90% deployment time reduction** via CI/CD
- Diagnosed N+1 query bottleneck; redesigned with **PostgreSQL indexing** achieving **sub-100ms** for 10,000+ records

PROJECTS

Pulumi – Contributor | Go, TypeScript, IaC

github.com/pulumi

- Submitted Go CLI features enabling multi-cloud (AWS/Azure/GCP) provisioning – under active review by core maintainers
- Analyzed **Raft/Paxos consensus** in distributed state synchronization layer – studied linearizability under concurrent operations and partial failures

SKILLS

ML & AI Infrastructure: vLLM, PagedAttention, Continuous Batching, PyTorch, LangChain, **LLM-as-a-judge evaluation**, LoRA/QLoRA

Languages: C++, Python, Go, TypeScript, Java, JavaScript, SQL, Bash, Rust

Distributed Systems: gRPC, Protocol Buffers, Circuit breakers, Exactly-once semantics, Fault tolerance, Kafka, Multi-AZ, Event-driven architecture

Cloud & Infrastructure: AWS (ECS Fargate, EKS, EC2, S3, RDS), Terraform, Kubernetes, Docker, CI/CD, CloudWatch, Prometheus, Grafana, GCP

Frameworks & Databases: Django, FastAPI, Node.js, React, REST APIs, PostgreSQL, Redis, MongoDB

AI Dev Tools: Claude Code, GitHub Copilot, Codex, Cursor