

Harry Nguyen

+1 470-667-9000 | nguyenharry2662@gmail.com | linkedin.com/in/harrynguyen26 | github.com/HarryNguyen2662

EDUCATION

Georgia State University

Expected May 2027

Bachelor of Science in Computer Science • GPA: 3.75/4.0

- Coursework: Distributed Systems, Operating Systems, Database Systems, Computer Networks, Data Structures
- Honors: Presidential Scholarship, Anita Andrew Fellowship, Roxie Alexander Memorial Award

EXPERIENCE

TiMoto AI | Software Engineer

Sep 2025 – Present

- Designed **gRPC inter-service layer** (Go/Python) with **exactly-once semantics** – traced production deadlock to shared resource acquisition conflicts; redesigned call sequencing; **100% evaluation success rate** at **sub-50ms p99**; **FastAPI + Django** REST layer serving external API clients; primary engineer on 3-person team
- Architected **multi-AZ AWS ECS Fargate** with **Terraform IaC**, circuit breaker pattern, and auto-rollback – **99.9% uptime SLO**; on-call rotation, root-cause analysis, runbooks, post-mortem; **44% infrastructure cost reduction** (\$40–60/mo)
- Built **Agentic AI evaluation pipeline** using **LangChain** orchestration + **LLM-as-a-judge** quality gates – automated deployment decisions via **Claude/GPT** APIs; replaced manual review with continuous enforcement; **zero production regressions**
- Deployed **vLLM with PagedAttention** for production LLM serving – KV cache optimization under concurrent load; continuous batching; **zero OOM failures**; **CloudWatch + Prometheus/Grafana** ML observability

Google | Software Engineer Intern – Chrome Browser

May 2025 – Aug 2025

- Designed **C++ IPC transport layer** with **Protocol Buffers** at **Google's security review standard** – evaluated Protobuf vs. custom serialization for memory safety, backwards-compatibility, and cross-language support; shipped to Chrome stable serving **3B+ active users** at **sub-50ms p99**; **95% test coverage**; zero production regressions
- Profiled settings navigation at 1,200ms p99; implemented **lock-free concurrent trie search** with **CAS operations** eliminating mutex contention – **96% latency reduction**, zero regressions; design adopted by senior Chrome engineers
- Architected event-driven **TypeScript/React** system with observer pattern – **68% feature delivery acceleration** across **25K+ lines** of Chromium at **95% test coverage**

Develop for Good | Software Engineer Intern

May 2024 – Aug 2024

- Designed stateless **BaaS on AWS** – selected **JWT** over session auth for stateless horizontal scalability; auto-scaling supporting **500+ concurrent users**; **90% deployment time reduction** via CI/CD (**GitHub Actions**)
- Diagnosed N+1 query bottleneck degrading API response to 3s+ on large datasets; redesigned with **PostgreSQL composite indexing** achieving **sub-100ms** for 10,000+ records

PROJECTS

Pulumi – Open Source Contributor | Go, TypeScript, Multi-Cloud IaC

github.com/pulumi

- Submitted **Go CLI features** for multi-cloud (AWS/Azure/GCP) provisioning to a **24.4K+ star** open-source project – developer tooling under active review by core maintainers; Go idiomatic patterns for CLI design and multi-provider abstraction
- Analyzed **Raft/Paxos consensus** in distributed state synchronization – studied linearizability guarantees and correctness under concurrent operations and partial failures

SKILLS

Distributed Systems: gRPC, Protocol Buffers, **Exactly-once semantics**, Circuit breakers, Multi-AZ, Fault tolerance, Event-driven architecture

Languages: Go, Python, C++, TypeScript, JavaScript, Java, Rust, **SQL**, Bash

Cloud & Infrastructure: AWS (ECS Fargate, EKS, EC2, S3, RDS), GCP, **Terraform**, Kubernetes, Docker, CI/CD (GitHub Actions), CloudWatch, Prometheus, Grafana

Frameworks & APIs: FastAPI, Django, React, Node.js, REST APIs, **PostgreSQL**, Redis, MongoDB, HTML/CSS

ML & AI Infrastructure: LangChain, LLM-as-a-judge evaluation, vLLM, PagedAttention, Generative AI, LLM APIs (Claude, GPT), PyTorch

AI Dev Tools: Claude Code, GitHub Copilot, Codex, Cursor