

Harry Nguyen

+1 470-667-9000 | nguyenharry2662@gmail.com | linkedin.com/in/harrynguyen26 | github.com/HarryNguyen2662

EDUCATION

Georgia State University

Expected May 2027

Bachelor of Science in Computer Science • **GPA: 3.75/4.0**

- Coursework: Distributed Systems, Operating Systems, Database Systems, Computer Networks, Data Structures
- Honors: Presidential Scholarship, Anita Andrew Fellowship, Roxie Alexander Memorial Award

EXPERIENCE

Google | Software Engineer Intern – Chrome Browser

May 2025 – Aug 2025

- Architected event-driven **TypeScript/React** system with observer pattern across **25K+ lines** of Chromium – **68% feature delivery acceleration** at **95% test coverage**; shipped to **3B+ active users**
- Profiled settings navigation at 1,200ms p99; designed **lock-free concurrent trie search** with **CAS operations** eliminating mutex contention – **96% latency reduction**, zero regressions; adopted by senior Chrome engineers
- Designed **C++ IPC transport layer** with **Protocol Buffers** – schema evolution, cross-language compatibility, **sub-50ms p99**; **95% test coverage**; zero production regressions

TiMoto AI | Software Engineer

Sep 2025 – Present

- Designed **FastAPI + Django** REST APIs and **gRPC** inter-service layer (**Go/Python**) – traced production deadlock to resource acquisition conflicts; redesigned call sequencing; **100% success rate** at **sub-50ms p99**; primary engineer on 3-person team
- Architected **multi-AZ AWS ECS Fargate + EKS** with **Terraform IaC** and auto-rollback – **99.9% uptime SLO**; **44% infrastructure cost reduction**; **Prometheus/Grafana + CloudWatch** observability stack
- Built **agentic AI pipeline** with **LangChain + LLM-as-a-judge** quality gates – autonomous quality gating for production deployments; zero production regressions; built with **Claude Code + Cursor**
- Deployed **vLLM with PagedAttention** for production LLM serving – KV cache optimization, continuous batching; **zero OOM failures** under concurrent load; **GitHub Actions** CI/CD with staged rollouts

Develop for Good | Software Engineer Intern

May 2024 – Aug 2024

- Diagnosed N+1 query bottleneck; redesigned with **PostgreSQL composite indexing** – **sub-100ms** API response for 10,000+ records; full **REST API** backend serving nonprofit client data
- Designed stateless **BaaS on AWS** – auto-scaling to **500+ concurrent users**, **90% deployment time reduction** via **GitHub Actions** CI/CD

PROJECTS

Pulumi – Open Source Contributor | *Go, TypeScript, Multi-Cloud IaC*

github.com/pulumi

- Submitted **Go CLI features** for multi-cloud (AWS/Azure/GCP) provisioning – open-source tooling under review by core maintainers; 24.4K+ star project
- Analyzed **Raft/Paxos consensus** in distributed state synchronization – linearizability guarantees and correctness under concurrent operations

SKILLS

Languages: TypeScript, C++, Python, JavaScript, Go, Java, Rust, **SQL**, Bash

Frameworks & APIs: React, Node.js, FastAPI, Django, **REST APIs**, gRPC, Protocol Buffers, **PostgreSQL**, Redis, MongoDB, HTML/CSS

Cloud & Infrastructure: AWS (ECS Fargate, EKS, EC2, S3, RDS), GCP, Terraform, K8s, Docker, CI/CD (GitHub Actions), CloudWatch, Prometheus, Grafana

Distributed Systems: Event-driven architecture, Circuit breakers, Exactly-once semantics, Multi-AZ, Fault tolerance, Async processing

ML & AI Infrastructure: LangChain, LLM-as-a-judge evaluation, vLLM, PagedAttention, LLM APIs (Claude, GPT), PyTorch

AI Dev Tools: Claude Code, Codex, Cursor, GitHub Copilot