Harry Nguyen

nguyenharry2662@gmail.com | linkedin.com/in/harrynguyen26 | timoto.ai | +1 470-667-9000

May 30, 2026

Glean Hiring Team
Glean Work, Inc.
Mountain View, CA

Re: Software Engineer, Billing & Revenue Platform

Glean is solving one of the hardest correctness problems in AI infrastructure: billing systems for workloads with unpredictable cost profiles. Exactly-once semantics, anomaly detection, and consumption metering under concurrency are not optional features—they are the foundation. I have built both layers in production.

At CoderPush, I designed idempotent billing APIs with exactly-once semantics under network partitions, backed by Redis distributed caching (85% cache hit rate at 9,000+ req/sec) and a DynamoDB hot-partition fix that improved read throughput by 30%. These are the same correctness guarantees billing systems require: no duplicate charges, no lost metering events, no data races under concurrent writes.

At TiMoto AI, I am the sole engineer behind an ML serving platform running on $40–60/month AWS infrastructure—a 44% cost reduction achieved by implementing vLLM PagedAttention guardrails and multi-AZ ECS Fargate with automated rollback via Terraform. Defining cost guardrails and anomaly thresholds for LLM workloads is part of my current day-to-day. When I resolved a production deadlock under concurrent gRPC evaluation requests—tracing the root cause to conflicting resource acquisition order and redesigning call sequencing—the result was 100% evaluation success rate at sub-50ms p99. Correctness under concurrency is not academic for me; it is a production constraint.

On top of these, I shipped C++ IPC changes at Google Chrome that reach 3 billion+ users (96% latency reduction on settings navigation), contributed Go CLI features to the Pulumi open-source project (24.4K stars), and hold a 3.75 GPA in Computer Science at Georgia State University with coursework in Distributed Systems, Operating Systems, and Database Systems.

Glean's consumption-based pricing infrastructure is exactly the kind of hard, consequential backend problem I want to own. The intersection of billing reliability and LLM cost structure is rare, and I have hands-on production experience on both sides. I would welcome the chance to show what I can build on your team.

Sincerely,

Harry Nguyen
nguyenharry2662@gmail.com
timoto.ai