

HARRY NGUYEN

+1 470-667-9000 | Atlanta, GA

✉ nguyenharry2662@gmail.com

in [linkedin.com/in/harrynguyen26](https://www.linkedin.com/in/harrynguyen26)

🐙 github.com/HarryNguyen2662

Education

Georgia State University

Atlanta, GA

Bachelor of Science in Computer Science, GPA: 3.75/4.0

Expected May 2027

- **Coursework:** Distributed Systems, Operating Systems, Database Systems, Computer Networks, Data Structures

Work Experience

TiMoto AI

Sep 2025 – Present

Software Engineer

Remote

- Primary engineer for backend, cloud infrastructure, and AI/ML serving on a 3-person team – designed, deployed, and operated distributed production systems end-to-end across the AI and infrastructure stack
- Architected vLLM inference engine with PagedAttention for LLM serving – eliminated KV cache memory fragmentation under concurrent load; deployed with continuous batching achieving zero OOM failures at production traffic
- Designed and maintained gRPC inter-service layer; resolved production deadlock under concurrent calls by tracing shared resource acquisition conflicts and redesigning call sequencing – 100% evaluation success rate at sub-50ms p99
- Architected multi-AZ ECS Fargate with Terraform IaC, CloudWatch observability, and circuit breaker pattern – 99.9% uptime, 44% cost reduction (\$40-60/month); auto-rollback on health check failure
- Maintained distributed production systems across gRPC, PostgreSQL, and ML serving layers – triaged and resolved incidents via root-cause analysis, documented runbooks, and applied post-mortem learnings; participated in on-call rotation

Google

May 2025 – Aug 2025

Software Engineer Intern – Chrome Browser

Mountain View, CA

- Designed C++ IPC transport layer with Protocol Buffers – selected Protocol Buffers over custom serialization for schema evolution and cross-language compatibility; shipped to Chrome stable channel serving 3B+ active users at sub-50ms p99, 10K+ req/sec
- Identified settings navigation as p99 bottleneck at 1,200ms; designed and implemented lock-free concurrent trie search eliminating mutex contention – 96% latency reduction, zero production regressions
- Architected event-driven TypeScript/React system with observer pattern decoupling UI state propagation – accelerated feature delivery 68% across 25K+ lines of Chromium at 95% test coverage
- Collaborated with Chrome infrastructure team to deliver design documents and code reviews at 95% test coverage; changes adopted into production branch reviewed by senior Chrome engineers

Develop for Good

May 2024 – Aug 2024

Software Engineer Intern

Remote

- Designed stateless BaaS on Amazon Web Services (AWS) – selected JWT over session auth for horizontal scalability; auto-scaling policy supporting 500+ concurrent users, 90% deployment time reduction via CI/CD (GitHub Actions)
- Diagnosed N+1 query bottleneck degrading response to 3s+ on large datasets; redesigned with PostgreSQL indexing achieving sub-100ms for 10,000+ records

Personal Projects

Pulumi – Contributor | Go, TypeScript, Infrastructure-as-Code (IaC)

- Submitted Go CLI features and bug fixes enabling multi-cloud (AWS/Azure/GCP) provisioning – under active review by core maintainers
- Analyzed Raft/Paxos consensus in distributed state synchronization layer – studied linearizability guarantees and correctness under concurrent operations and partial failures

Technical Skills

Distributed Systems: gRPC, Protocol Buffers, Circuit breakers, Exactly-once semantics, Fault tolerance, Kafka, Multi-AZ, Event-driven architecture

ML & AI Infrastructure: vLLM, PagedAttention, Continuous Batching, LoRA/QLoRA, PyTorch, LangChain, LLM-as-a-judge evaluation

Languages: C++, Python, Go, TypeScript, Java, JavaScript, SQL, Bash, Rust

Cloud & Infrastructure: AWS (ECS Fargate, EKS, EC2, S3, RDS), Terraform, Kubernetes, Docker, CI/CD, CloudWatch, Prometheus, Grafana, GCP

Frameworks & Databases: Django, FastAPI, Node.js, React, REST APIs, PostgreSQL, Redis, MongoDB

AI Dev Tools: Claude Code, GitHub Copilot, Codex, Cursor