# Harry Nguyen

Atlanta, GA | +1 470-667-9000 | nguyenharry2662@gmail.com | linkedin.com/in/harrynguyen26 | github.com/HarryNguyen2662

## Education

**Georgia State University** **Atlanta, GA**

*Bachelor of Science in Computer Science, GPA: 3.75/4.0* *Expected May 2027*

- **Coursework:** Distributed Systems, Operating Systems, Database Systems, Computer Networks, Data Structures

## Work Experience

**TiMoto AI** **Sep 2025 – Present**

*Software Engineer*

- Restored **100% evaluation success rate** at **sub-50ms p99** by resolving a production **gRPC** deadlock – replaced inconsistent lock-acquisition ordering under concurrent calls with a single enforced sequencing rule across the inter-service layer
- Eliminated **OOM failures** under concurrent production load by replacing naive batching with a **vLLM/PagedAttention** inference engine, solving KV cache memory fragmentation for LLM serving
- Prevented single-zone failures from taking production down, sustaining **99.9% uptime** and cutting infrastructure cost **44%** ($40-60/month), by architecting a multi-AZ ECS Fargate deployment with **Terraform IaC**, CloudWatch observability, and an auto-rollback circuit breaker pattern

**Google** **May 2025 – Aug 2025**

*Software Engineer Intern – Chrome Browser*

- Turned a visibly sluggish settings search into an instant one for power users – **96%** p99 latency cut (1,200ms down to sub-50ms), zero production regressions – by replacing a mutex-contended search path with a **lock-free concurrent trie**
- Shipped a **C++ IPC transport layer** to Chrome's stable channel – serving **3B+ active users** at **sub-50ms p99** and **10K+ req/sec** – by choosing Protocol Buffers over custom serialization for schema evolution and cross-language compatibility
- Landed production-ready changes into Chrome's codebase – reviewed and adopted by senior engineers – by delivering design docs and code reviews at **95% test coverage** alongside the Chrome infrastructure team

**Develop for Good** **May 2024 – Aug 2024**

*Software Engineer Intern*

- Cut API response time from 3+ seconds to **sub-100ms** on **10,000+ record** datasets by diagnosing an N+1 query bottleneck and redesigning the data layer with **PostgreSQL indexing**
- Scaled a backend to support **500+ concurrent users** and cut deployment time **90%** by designing a stateless AWS BaaS with **Python**, JWT auth, and a GitHub Actions CI/CD auto-scaling policy

## Personal Projects

**Pulumi – Contributor | *Go, TypeScript, Infrastructure-as-Code (IaC)***

- Submitted Go CLI features and bug fixes enabling multi-cloud (AWS/Azure/GCP) provisioning – under active review by core maintainers
- Analyzed **Raft/Paxos consensus** in distributed state synchronization layer – studied linearizability guarantees and correctness under concurrent operations and partial failures

## Technical Skills

**Languages**: C++, Python, Go, TypeScript, Java, JavaScript, SQL, Bash, Rust
**Distributed Systems**: gRPC, Protocol Buffers, Circuit breakers, Exactly-once semantics, Fault tolerance, Kafka, Multi-AZ, Event-driven architecture
**Cloud & Infrastructure**: AWS (ECS Fargate, EKS, EC2, S3, RDS), Terraform, Kubernetes, Docker, CI/CD, CloudWatch, Prometheus, Grafana, GCP
**ML & AI Infrastructure**: vLLM, PagedAttention, Continuous Batching, LoRA/QLoRA, PyTorch, LangChain, LLM-as-a-judge evaluation
**Frameworks & Databases**: Django, FastAPI, Node.js, React, REST APIs, PostgreSQL, Redis, MongoDB
**AI Dev Tools**: Claude Code, GitHub Copilot, Codex, Cursor